

# Apache Airflow Workflow Orchestration



# What is Airflow?

## Definition of Airflow

Apache Airflow is an open-source tool designed to programmatically author, schedule, and monitor workflows as directed acyclic graphs (DAGs).

## Purpose and Use Cases

Airflow automates complex data pipelines, enabling reliable execution, scheduling, and monitoring of workflows in data engineering, ETL, machine learning, and DevOps.

## Why Airflow for Workflow Orchestration

Airflow offers dynamic pipeline generation, extensibility with custom operators, robust scheduling, and a rich UI for monitoring and managing workflows.

## Key Features

Features include DAG-based workflow management, modular architecture, rich user interface, extensive operator support, and integration with many platforms.


## Common Use Cases

Used for ETL processes, data warehousing, machine learning workflows, infrastructure automation, and batch job orchestration in various industries.



# Airflow Architecture & Components

## HOW APACHE AIRFLOW WORKS



# Airflow Architecture & Components

---

## Scheduler

The Scheduler triggers tasks based on defined DAG schedules, deciding when and what tasks to run next to ensure timely workflow execution.

## Executor

The Executor handles task execution by distributing work to workers; it supports different modes like Local, Sequential, Celery, and Kubernetes Executors.

## Worker

Workers are the machines or processes that perform the actual task execution assigned by the Executor, running task instances in parallel.

## Web Server

The Webserver hosts the Airflow UI, allowing users to monitor DAGs, visualize workflows, trigger tasks, and view logs through a graphical interface.

## Metadata Database

Stores all Airflow metadata including DAG definitions, task states, and scheduling information, serving as the single source of truth.

# DAGs and Task Lifecycle

## DAGs Concept

A Directed Acyclic Graph (DAG) represents a workflow as a collection of tasks with dependencies, ensuring tasks execute in a defined order without cycles.

## Task Lifecycle Stages

Tasks transition through states like **queued**, **running**, **success**, **failed**, and **skipped**, reflecting their progress and outcomes during execution.

## Relationship Between Tasks and DAGs

Tasks are the nodes within a DAG; the DAG defines how tasks depend on each other and their execution sequence.

## DAG Structure Overview

A DAG is defined in Python code, specifying tasks, dependencies, scheduling intervals, and other execution parameters.

## How Tasks Are Executed

The scheduler triggers tasks based on dependencies and schedule, the executor runs the tasks, and workers process them in parallel or sequentially.



# Airflow Installation & Prerequisites

## Prerequisites for Airflow

Ensure **Python 3.7+** is installed along with pip. Additional system dependencies like database drivers and build tools may be required depending on Airflow features.

## Installing via pip

Run '**pip install apache-airflow**' with optional extras like 'postgres', 'celery' according to your setup. It offers flexibility for different backend integrations.

## Installing via conda

Use '**conda install -c conda-forge airflow**' for an isolated environment. Conda handles dependencies well, making it ideal for data science workflows.

## Installing via Docker

Use official Airflow Docker images to avoid complex setups. Docker ensures consistency across environments and simplifies scaling and deployment.

## Config File Overview (airflow.cfg)

The **airflow.cfg** file contains key configurations like **executor type**, **database connection**, **logging**, and **scheduler settings**. Customize it to optimize Airflow behavior for your environment.



# Airflow File System Overview



## DAGs Folder

This folder stores all Directed Acyclic Graph (DAG) definitions as Python scripts, which Airflow scans to schedule and execute workflows.

## Logs Folder

Contains execution logs for tasks and **DAG runs**, helping users **monitor progress, debug issues**, and **audit workflow** performance.

## Plugins

Custom **operators, sensors, hooks**, and **interfaces** are placed here to extend Airflow's capabilities beyond its built-in functions.

## Config Files

Includes the airflow.cfg configuration file that defines system parameters, environment settings, and executor options for Airflow.

# Airflow User Interface Overview

## DAG Dashboard

Central hub displaying all DAGs with their **status, schedule, and recent run details** for quick overview and access.

## Graph View

Visual representation of DAG tasks and dependencies as a directed acyclic graph, helping users understand execution flow.

## Tree View

Hierarchical view showing task status over time, useful for tracking progress and identifying failed or skipped tasks.

## Task Monitoring

Real-time updates on individual task execution with logs, duration, and status to facilitate troubleshooting and performance tracking.

## Triggering DAGs

Manual and automated options to start DAG runs directly from the UI, enabling control over workflow execution timing.



# Common Airflow CLI Commands

## `airflow db init`

Initializes the **Airflow metadata database**, creating all necessary tables and structures for the system to operate.

## `airflow scheduler`

Starts the scheduler process that triggers task execution following the defined DAG schedules and dependencies.

## `airflow webserver`

Launches the Airflow web interface, providing a **dashboard to monitor, manage, and trigger DAGs visually**.

## `airflow dags list`

Lists all DAGs currently available in the Airflow environment, allowing users to see which workflows are defined.

## `airflow tasks run`

Manually triggers the execution of a specific task within a DAG, useful for testing or rerunning tasks on demand.



# Types of Airflow Operators

## PythonOperator

Executes Python functions or callables within a DAG, enabling integration of complex logic and data processing directly in workflows.

## BashOperator

Runs bash commands or shell scripts, allowing system-level command execution as part of the workflow automation.

## DummyOperator

Acts as a placeholder or no-op task, useful for structuring DAGs or marking task dependencies without performing actions.

## SensorOperator

Waits for a specific condition or external event to be met before proceeding, such as file availability or database updates.

## Custom Operators

Users can create their own operators by extending base classes to handle unique or complex tasks tailored to specific workflows.

# Airflow Executors Explained

---



## LocalExecutor

Runs tasks in parallel on the same machine using multiprocessing, suitable for moderate workloads and local setups.



## SequentialExecutor

Executes tasks one at a time sequentially, ideal for development and testing but not for production use due to limited concurrency.



## CeleryExecutor

Distributes task execution across multiple worker nodes using Celery and message brokers, enabling scalable and distributed workflows.



## KubernetesExecutor

Launches each task in its own Kubernetes pod, providing dynamic scaling, isolation, and cloud-native orchestration benefits.

---

## Data Passing and Variables in Airflow

---



### What are XComs?

XComs (Cross-communications) are a built-in feature that allows tasks to exchange small pieces of data asynchronously during DAG execution.

### Passing Data Between Tasks

Tasks use XComs to push and pull data, enabling downstream tasks to access outputs from upstream tasks within the same DAG run.

### Airflow Variables

Variables are **key-value pairs** stored centrally in Airflow's **metadata** database, used for storing configuration or constant values accessible by all tasks.

### Example with PythonOperator

A PythonOperator can push data to XComs using 'xcom\_push' and pull it in downstream tasks with 'xcom\_pull', facilitating dynamic data-driven workflows.

---

## Best Practices & Debugging Tips

### Best Practices in DAG Design

Design DAGs with clear task dependencies and modular, reusable tasks. Use meaningful task IDs, avoid hardcoding values, and leverage Airflow variables and connections for flexibility.

### Debugging Tips

Monitor logs and task states via the Airflow UI. Use verbose logging for detailed error analysis. Test tasks independently and apply retries and alerts for failure handling.

### Real-world Case Studies

Examples include automating ETL pipelines, orchestrating machine learning workflows, and managing daily batch jobs, which demonstrate improved reliability and scalability with Airflow.

---

**Thank you**

---

